

Case Study: 2 or 7?

Machine Learning - Section 3.1.5

Marc Omar Haddad

Published: 23 February, 2020

Updated: 24 February, 2020

In the simple examples we've examined up to now, we only had one predictor. We actually do not consider these machine learning challenges, which are characterized by having many predictors.

— Dr. Rafael Irizarry

There are 784 predictors/features in a digit-reading machine learning algorithm. But for this case study we will select a subset of the aforementioned predictors. This subcategory will consist of **2 predictors** and **2 categories**.

Our goal for this lesson: Build an algorithm to determine if a digit is 2 or 7 (our 2 categories) from predictors. Our 2 predictors will be:

1. **The proportion of dark pixels in upper-left quadrant X_1 .**
2. **The proportion of dark pixels in lower-right quadrant X_2 .**

We will use a random sample of 1,000 digits from the 60,000 digits in our training set. 500 digits will be allocated to our custom **training set** and 500 digits will be allocated to our custom **test set**. These sets can be accessed within the `dslabs` package by running `data("mnist_27")`.

First, we plot our two predictors from the data to explore it. We will use colors to denote labels.

```
data("mnist_27")

data = mnist_27$train %>% drop_na()
data %>% ggplot(aes(x_1, x_2, color = y)) + geom_point()
```



We can clearly see some patterns in the above plot:

- 7's are associated with large black values in the upper-left quadrant (approx. $X_1 \geq 0.3$).
- 2's are associated with mid-range black values in the lower-right quadrant (approx. $0.2 < X_2 < 0.35$).

In this case study, we will attempt to fit a logistic regression model to our data.

Mathematically Expressing Logistic Regression

We will assume that “7” is the “positive” outcome (i.e. $Y = 1$).

Our model can be expressed mathematically like so:

$$\begin{aligned} g\{p(x_1, x_2)\} &= g\{\Pr(Y = 1 \mid X_1 = x_1, X_2 = x_2)\} \\ &= \beta_0 + \beta_1 x_1 + \beta_2 x_2 \end{aligned}$$

Or, alternatively:

$$\begin{aligned} p(x_1, x_2) &= \Pr(Y = 1 \mid X_1 = x_1, X_2 = x_2) \\ &= g^{-1}(\beta_0 + \beta_1 x_1 + \beta_2 x_2) \end{aligned}$$

With $g^{-1}(x)$ being the inverse of the logistic function $g(x)$:

$$g(x) = \log \frac{x}{1-x}$$

$$g^{-1}(x) = \frac{\exp(x)}{1 + \exp(x)}$$

In simple english,

First method:

- The **logistic transformation** ($g\{\}$) of the **probability** of the **outcome being “7”** ($Y = 1$) *given*:
 - The black value in the upper-left quadrant (\mathbf{x}_1) *and*
 - The black value in the lower-right quadrant (\mathbf{x}_2)
- Is *equal to* the **linear function** $\beta_0 + \beta_1 x_1 + \beta_2 x_2$

Alternative method:

- The **probability** of the **outcome being “7”** ($Y = 1$) *given*:
 - The black value in the upper-left quadrant (\mathbf{x}_1) *and*
 - The black value in the lower-right quadrant (\mathbf{x}_2)
- Is *equal to* the **inverse logistic transformation** ($g^{-1}\{\}$) of the **linear function** $\beta_0 + \beta_1 x_1 + \beta_2 x_2$

Fitting and Applying Logistic Regression

Despite the slight challenge involved in mathematically expressing logistic regression, fitting our model in R is fairly straight forward:

```
fit = glm(y ~ x_1 + x_2, data = mnist_27$train, family = "binomial")
```

With our fit we can now build a **decision rule** based on the **estimate of the conditional probability**.

Decision Rule: If the estimate of the conditional probability is greater than 0.5, predict “7”.

```
p_hat = predict(fit, newdata = mnist_27$test)
y_hat = ifelse(p_hat > 0.5, 7, 2) %>% factor()

# Note: Our matrix indicates "2" as our positive class, but that's okay.
confusionMatrix(y_hat, mnist_27$test$y)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  2  7
##           2 92 34
##           7 14 60
##
##           Accuracy : 0.76
##           95% CI : (0.6947, 0.8174)
##       No Information Rate : 0.53
##       P-Value [Acc > NIR] : 1.668e-11
##
##           Kappa : 0.5124
##
## Mcnemar's Test P-Value : 0.006099
##
##           Sensitivity : 0.8679
##           Specificity : 0.6383
##       Pos Pred Value : 0.7302
##       Neg Pred Value : 0.8108
##           Prevalence : 0.5300
##       Detection Rate : 0.4600
##       Detection Prevalence : 0.6300
##       Balanced Accuracy : 0.7531
##
##       'Positive' Class : 2
##
```

As we can see in the `confusionMatrix()` results, we obtained an accuracy of 0.76.